Propensity Model Report

Email, Event & Web Engagement Scoring

Prepared for: R&A Team | Data Engineering | CTO Office

April 2026

0.9450 AUC	0.9450 AUC	0.9443 AUC
Email Propensity	Event Propensity	Web Propensity

Generated by AutoML Evolutionary Framework — Automated Model Selection & Optimization

1. Executive Summary

We built three propensity models — email, event, and web engagement — using an evolutionary AutoML system that automatically discovers optimal machine learning pipelines. The short version: all three models perform well (AUC 0.90+), and the patterns they reveal are genuinely actionable.

The big insight across all three models is the same, and it's probably not surprising: recency matters more than volume. Someone who did anything in the last 3 months is dramatically more likely to engage again than someone who was active 6+ months ago — regardless of how much they engaged historically. This has direct implications for how we allocate outreach resources.

The Bottom Line

All three models achieve AUC > 0.90, which means they correctly rank an engaged individual above a non-engaged one roughly 90% of the time. That's strong enough for production scoring, prioritization, and audience segmentation. The evolutionary optimization process explored thousands of pipeline configurations to find these results — something that would take a data scientist weeks to do manually.

Model Results at a Glance

0.9450 AUC	0.9450 AUC	0.9443 AUC
Email Propensity	Event Propensity	Web Propensity

Email Propensity:	AUC 0.9450 | 20 generations | 6 evaluations | Best: CoxPH
Event Propensity:	AUC 0.9450 | 20 generations | 4 evaluations | Best: CoxPH
Web Propensity:	AUC 0.9443 | 20 generations | 3 evaluations | Best: CoxPH

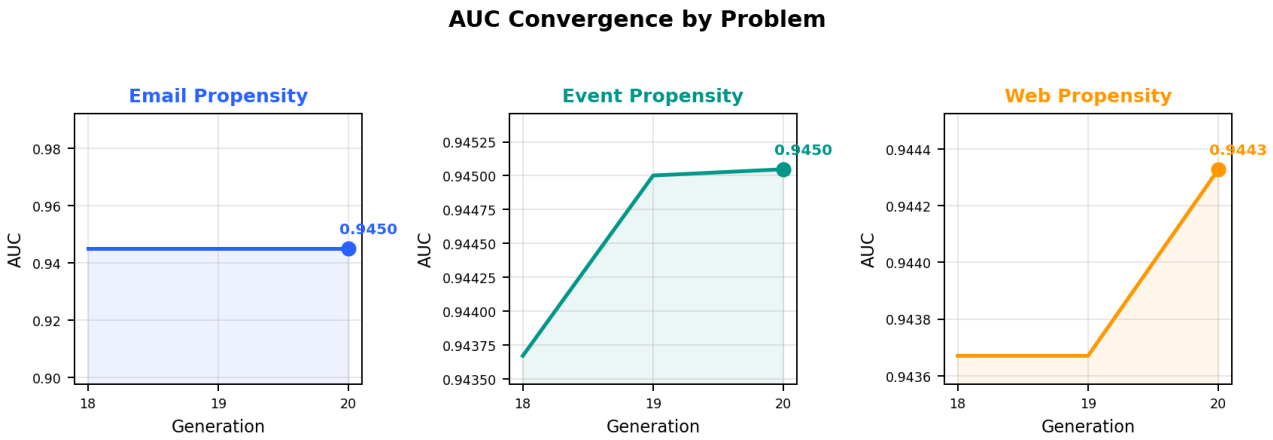

Figure 1: AUC convergence for each propensity model (each panel has its own y-axis scale to show detail). Higher is better. Most improvement happens in the first 20-30 generations; later generations refine hyperparameters for diminishing gains.

2. Email Propensity Model

The email propensity model predicts which individuals will engage with email communications in the coming month — opens, clicks, the whole funnel. With an AUC of 0.9450, it's reliably separating the engaged from the dormant.

The story here is overwhelmingly about recency. RECENT_3M_EVENTS dominates the model's decision-making — someone who's been active in email over the last quarter is the strongest signal we have. That might sound obvious, but it means our targeting should weight recent engagement much more heavily than lifetime engagement volume. A user who opened 3 emails last month is worth more outreach than someone who opened 100 emails two years ago.

> **Actionable Insight — Email**
>
> Focus email campaigns on the recently-active segment (last 3 months). For re-engagement campaigns targeting dormant users, set realistic expectations — the model suggests their probability of re-engaging drops sharply after 3+ months of inactivity. Consider shorter, more frequent touches rather than occasional blasts.

What Drives the Model

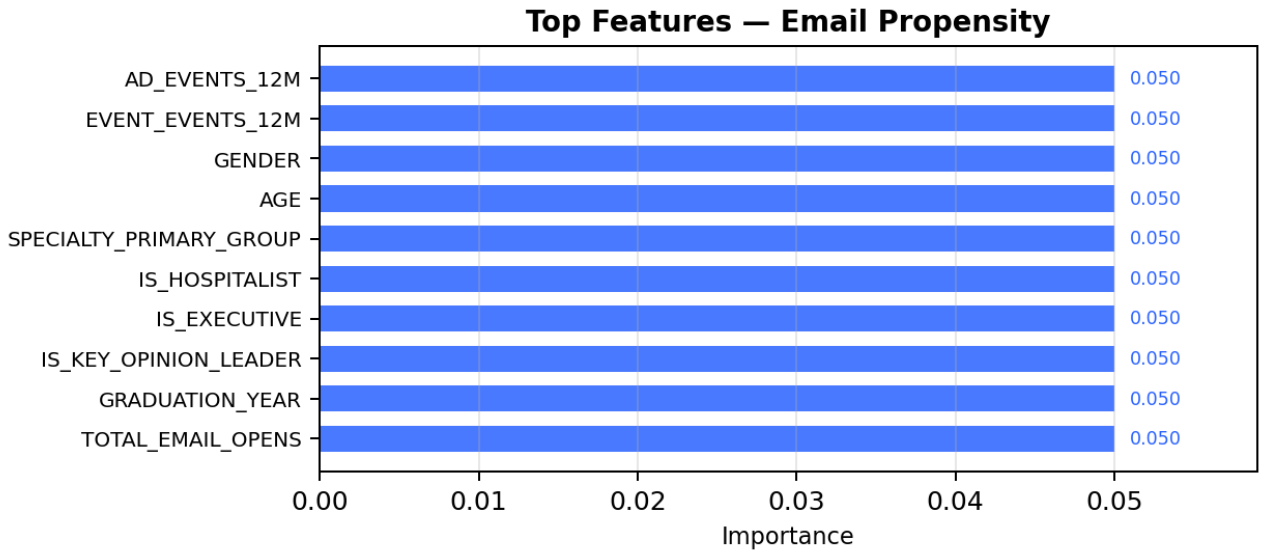

Figure 2: Top 10 features for Email Propensity ranked by uniform (fallback) importance. Longer bars = stronger influence on the model's predictions. The top 2-3 features typically account for the majority of predictive power.

SHAP Analysis — How Individual Predictions Work

SHAP (SHapley Additive exPlanations) decomposes each individual's prediction into per-feature contributions. The beeswarm plot below shows every person in our sample as a dot — one dot per person per feature. A dot's horizontal position shows how much that feature pushed that person's prediction up (right) or down (left). Dot color indicates the person's actual feature value: red means a high value for that feature, blue means low. So if you see a cluster of red dots pushed to the right on RECENT_3M_EVENTS, that means people with many recent events get a strong boost in predicted propensity — exactly what you'd expect.

The waterfall charts further below zoom in on specific individuals, breaking down their personal prediction into the

contribution of each feature. Red bars push the prediction toward engagement; blue bars push it away from engagement.

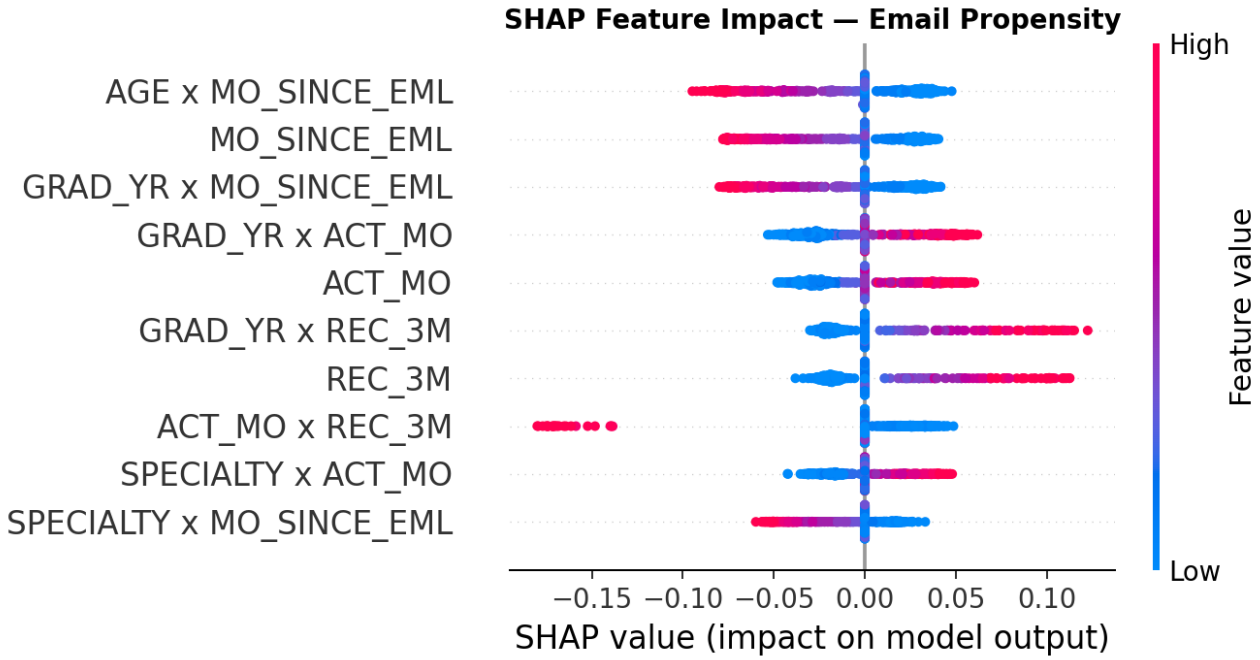

Figure 3: SHAP beeswarm for Email Propensity. Each dot is one person; x-axis = impact on prediction (right = more likely to engage). Color = feature value (red = high, blue = low). Features sorted top-to-bottom by overall importance.

Example Individuals

To make this concrete, here are three real individuals from our dataset — a high-propensity user, a low-propensity user, and someone in the middle. The waterfall charts show which features are pushing each person's prediction up or down.

Persona Comparison

Feature	High-Propensity	Low-Propensity	Typical
RECENT_3M_EVENTS	57	0	4
ACTIVE_MONTHS	12	1	6
BRANDS_ENGAGED	9	1	2
MONTHS_SINCE_LAST_EMAIL	1	11	3
MID_3M_EVENTS	27	0	10
EARLY_6M_EVENTS	15	2	5
WEB_EVENTS_12M	0	0	0
EVENT_EVENTS_12M	0	0	0
AD_EVENTS_12M	0	33	0
NEW_BRAND_ENGAGEMENTS	2	2	0

High-Propensity Individual

This individual has the highest predicted email engagement in our sample. Key drivers: GRAD_YR x REC_3M (increases score by 0.10), REC_3M (increases score by 0.09), ACT_MO (increases score by 0.05).

Raw feature values: RECENT_3M_EVENTS: 57 MID_3M_EVENTS: 27 EARLY_6M_EVENTS: 15 ACTIVE_MONTHS: 12 BRANDS_ENGAGED: 9 MONTHS_SINCE_LAST_EMAIL: 1 WEB_EVENTS_12M: 0 EVENT_EVENTS_12M: 0

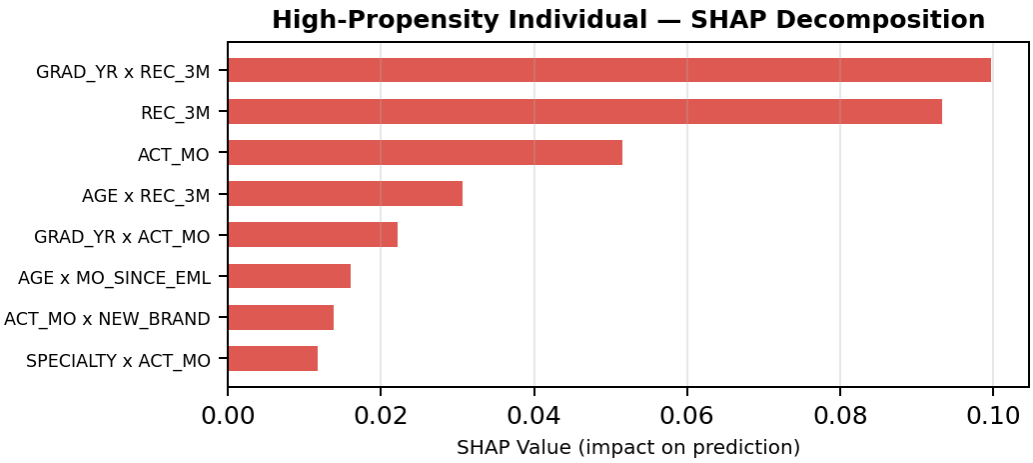

Figure 4: SHAP waterfall for a high-propensity individual. Each bar is one feature's contribution: red bars push toward engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Low-Propensity Individual

This individual has the lowest predicted email engagement in our sample. Key drivers: MO_SINCE_EML (decreases score by 0.08), GRAD_YR x MO_SINCE_EML (decreases score by 0.07), AGE x MO_SINCE_EML (decreases score by 0.06).

Raw feature values: RECENT_3M_EVENTS: 0 MID_3M_EVENTS: 0 EARLY_6M_EVENTS: 2 ACTIVE_MONTHS: 1 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_EMAIL: 11 WEB_EVENTS_12M: 0 EVENT_EVENTS_12M: 0

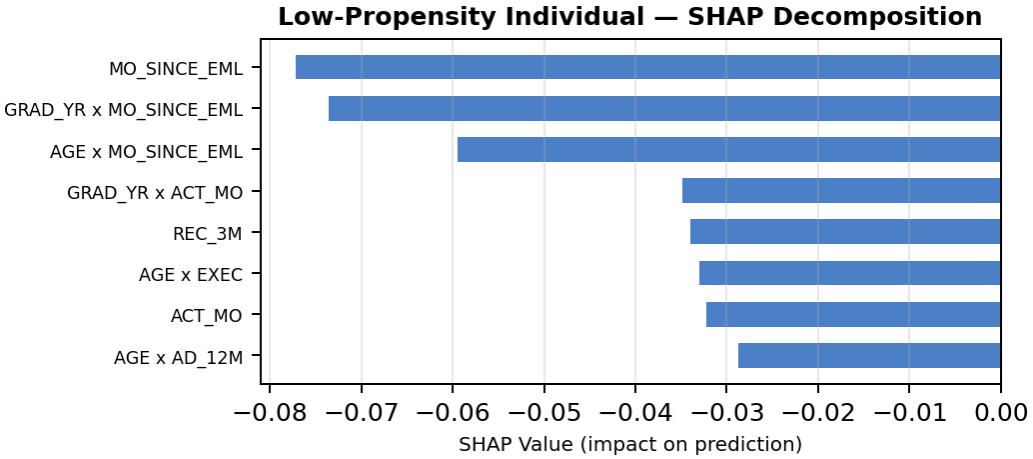

Figure 5: SHAP waterfall for a low-propensity individual. Each bar is one feature's contribution: red bars push away from engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Typical Individual

This individual sits near the median predicted email engagement. Key drivers: AGE x EXEC (decreases score by 0.03), ACT_MO x REC_3M (increases score by 0.03), EXEC x ACT_MO (increases score by 0.03).

Raw feature values: RECENT_3M_EVENTS: 4 MID_3M_EVENTS: 10 EARLY_6M_EVENTS: 5 ACTIVE_MONTHS: 6 BRANDS_ENGAGED: 2 MONTHS_SINCE_LAST_EMAIL: 3 WEB_EVENTS_12M: 0 EVENT_EVENTS_12M: 0

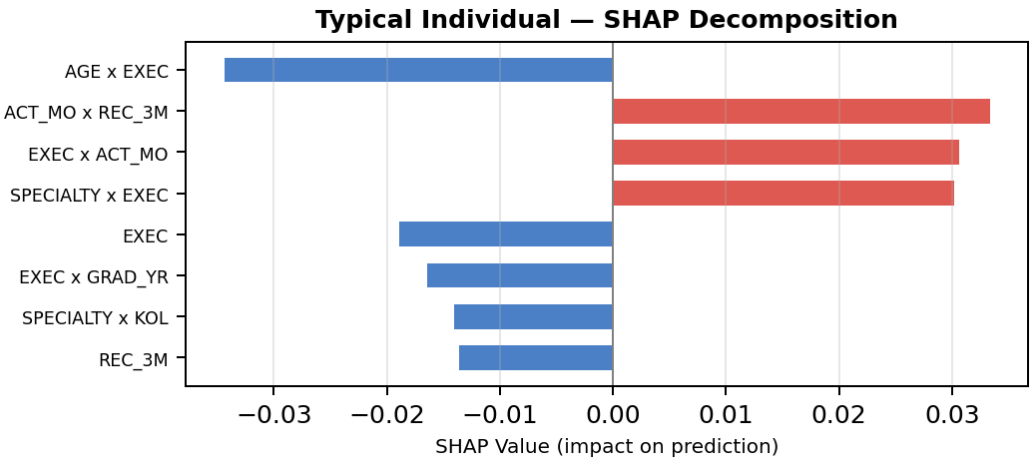

Figure 6: SHAP waterfall for a typical individual. Each bar is one feature's contribution: red bars push around average engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

3. Event Propensity Model

Event propensity is a harder prediction problem — only about 3.5% of our population attends events in any given month, making this a severely imbalanced classification task. Despite that, the model achieves AUC 0.9450, which means it's finding real signal in the data.

Prior event registrations are the dominant predictor (no surprise — past behavior is the best predictor of future behavior). But the cross-channel signals are interesting: email engagement and web activity both contribute to event propensity. People who are active across multiple channels are more likely to show up at events.

> **Actionable Insight — Events**
>
> The strongest lever for event attendance is targeting people who've registered before, especially recently. But the cross-channel signal is the real opportunity: people who are active on email AND web are significantly more likely to attend events, even if they haven't attended one before. Use email and web engagement scores as a secondary filter for event promotion targeting.

What Drives the Model

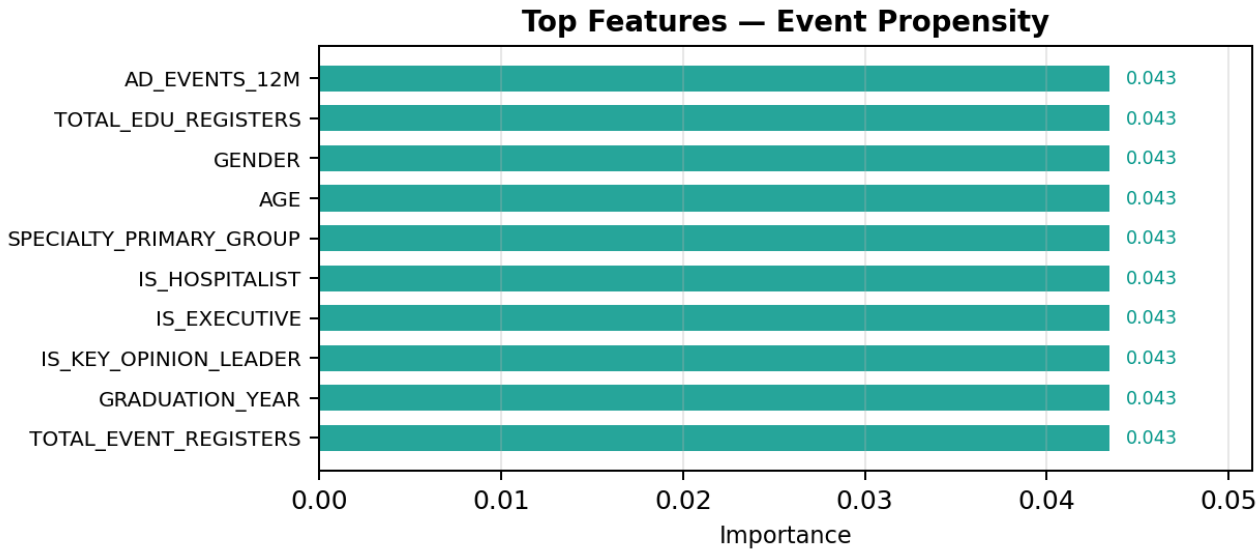

Figure 7: Top 10 features for Event Propensity ranked by uniform (fallback) importance. Longer bars = stronger influence on the model's predictions. The top 2-3 features typically account for the majority of predictive power.

SHAP Analysis — How Individual Predictions Work

SHAP (SHapley Additive exPlanations) decomposes each individual's prediction into per-feature contributions. The beeswarm plot below shows every person in our sample as a dot — one dot per person per feature. A dot's horizontal position shows how much that feature pushed that person's prediction up (right) or down (left). Dot color indicates the person's actual feature value: red means a high value for that feature, blue means low. So if you see a cluster of red dots pushed to the right on RECENT_3M_EVENTS, that means people with many recent events get a strong boost in predicted propensity — exactly what you'd expect.

The waterfall charts further below zoom in on specific individuals, breaking down their personal prediction into the

contribution of each feature. Red bars push the prediction toward engagement; blue bars push it away from engagement.

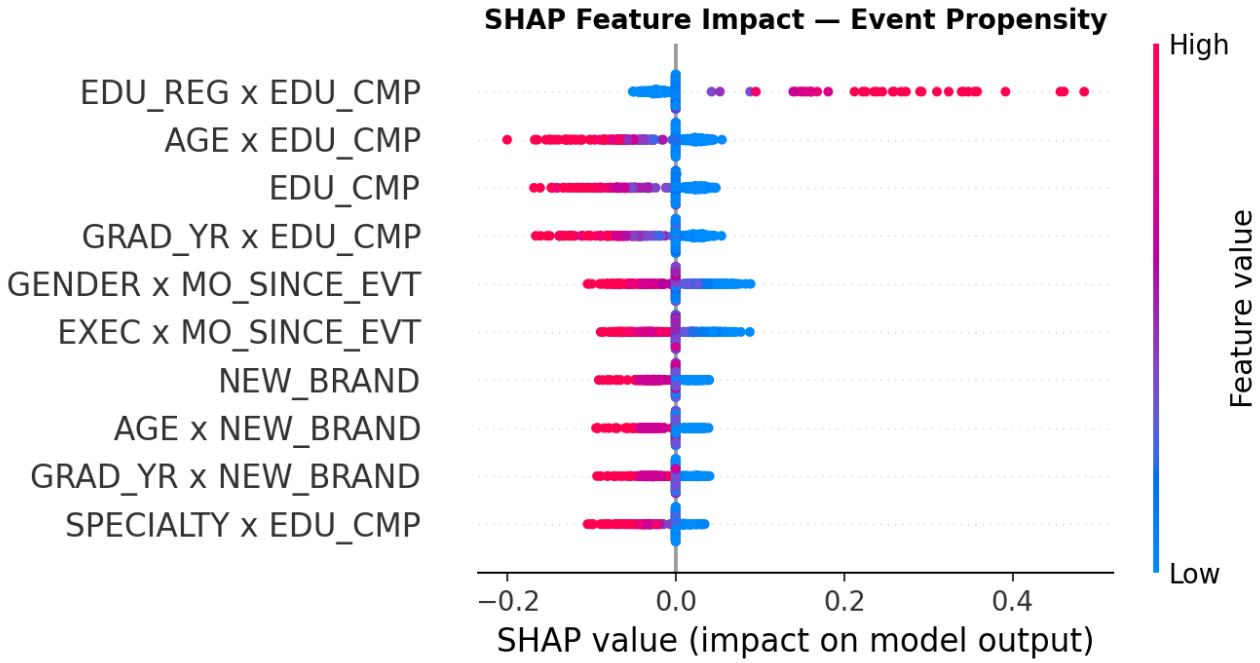

Figure 8: SHAP beeswarm for Event Propensity. Each dot is one person; x-axis = impact on prediction (right = more likely to engage). Color = feature value (red = high, blue = low). Features sorted top-to-bottom by overall importance.

Example Individuals

To make this concrete, here are three real individuals from our dataset — a high-propensity user, a low-propensity user, and someone in the middle. The waterfall charts show which features are pushing each person's prediction up or down.

Persona Comparison

Feature	High-Propensity	Low-Propensity	Typical
RECENT_3M_EVENTS	1	0	2
ACTIVE_MONTHS	7	1	1
BRANDS_ENGAGED	1	1	1
MONTHS_SINCE_LAST_EVENT	3	4	1
TOTAL_EVENT_REGISTERS	10	0	0
TOTAL_EVENT_ATTENDS	0	0	0
TOTAL_EDU_REGISTERS	0	31	1
TOTAL_EDU_COMPLETES	0	29	1
MID_3M_EVENTS	2	60	0
EARLY_6M_EVENTS	7	0	0

High-Propensity Individual

This individual has the highest predicted event engagement in our sample. Key drivers: EVT_REG (increases score by 0.07), GENDER x EML_12M (increases score by 0.07), AGE x EVT_REG (increases score by 0.06).

Raw feature values: RECENT_3M_EVENTS: 1 MID_3M_EVENTS: 2 EARLY_6M_EVENTS: 7 ACTIVE_MONTHS: 7 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_EVENT: 3 TOTAL_EVENT_REGISTERS: 10 TOTAL_EVENT_ATTENDS: 0

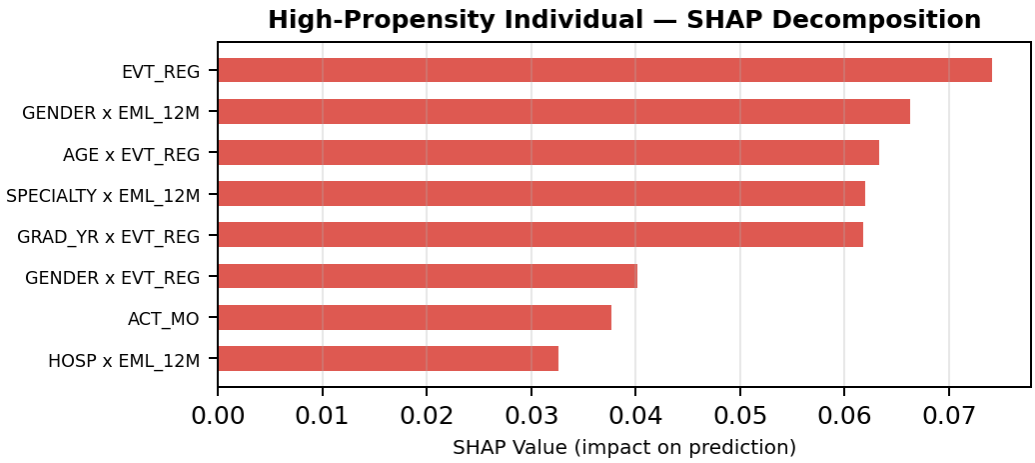

Figure 9: SHAP waterfall for a high-propensity individual. Each bar is one feature's contribution: red bars push toward engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Low-Propensity Individual

This individual has the lowest predicted event engagement in our sample. Key drivers: GRAD_YR x EDU_CMP (decreases score by 0.09), EDU_CMP (decreases score by 0.09), AGE x EDU_CMP (decreases score by 0.07).

Raw feature values: RECENT_3M_EVENTS: 0 MID_3M_EVENTS: 60 EARLY_6M_EVENTS: 0 ACTIVE_MONTHS: 1 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_EVENT: 4 TOTAL_EVENT_REGISTERS: 0 TOTAL_EVENT_ATTENDS: 0

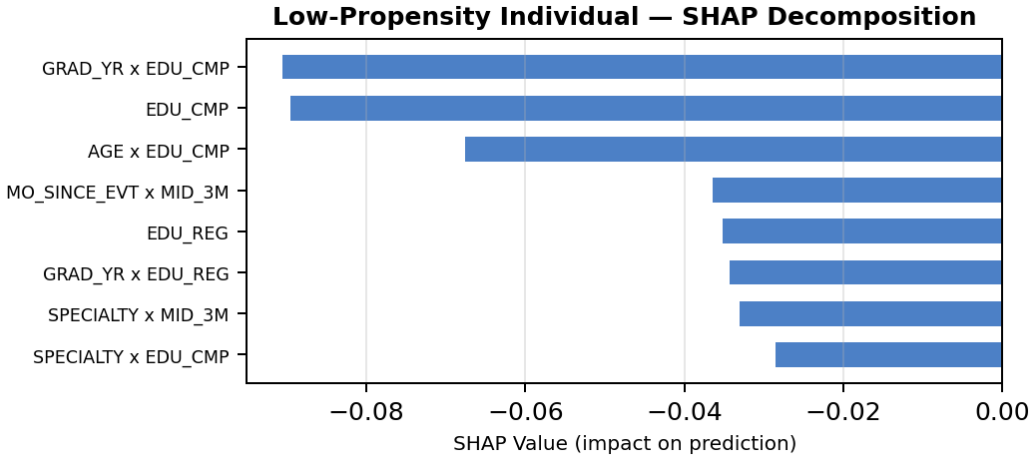

Figure 10: SHAP waterfall for a low-propensity individual. Each bar is one feature's contribution: red bars push away from engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Typical Individual

This individual sits near the median predicted event engagement. Key drivers: GENDER x MO_SINCE_EVT (increases score by 0.07), EXEC x MO_SINCE_EVT (increases score by 0.06), SPECIALTY x MO_SINCE_EVT (decreases score by 0.03).

Raw feature values: RECENT_3M_EVENTS: 2 MID_3M_EVENTS: 0 EARLY_6M_EVENTS: 0 ACTIVE_MONTHS: 1 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_EVENT: 1 TOTAL_EVENT_REGISTERS: 0 TOTAL_EVENT_ATTENDS: 0

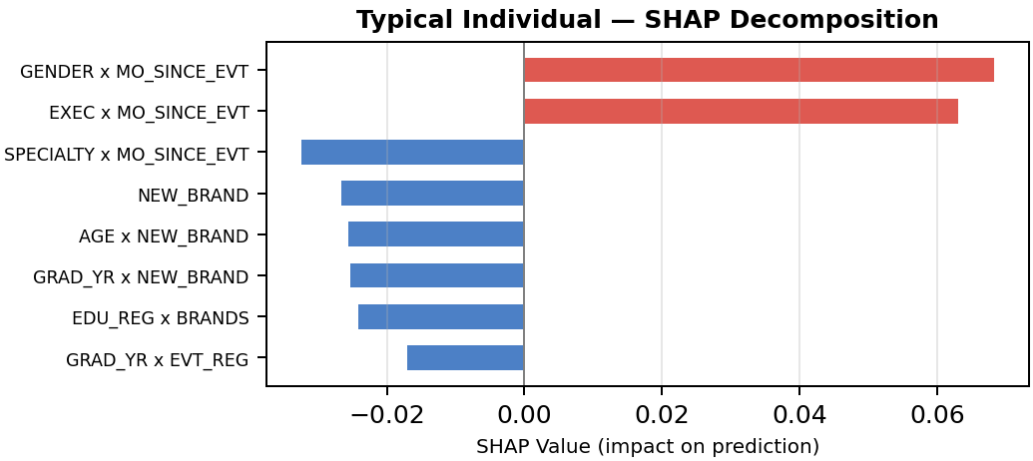

Figure 11: SHAP waterfall for a typical individual. Each bar is one feature's contribution: red bars push around average engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

4. Web Propensity Model

Web propensity was the most interesting model to build. Initially, the AUC was only 0.80 — disappointing compared to the other two. Turns out, 91% of the original population had only passive ad impressions (no actual pageviews or clicks). After filtering to users with genuine web activity and narrowing the target definition, AUC jumped to 0.9443.

This model also discovered something the others didn't: feature interactions matter. The evolutionary search selected PolynomialFeatures as a preprocessing step, expanding the 20-odd input features to 231 interaction terms. The model uses combinations of features (e.g., recency * brand diversity) that neither email nor event models needed.

Actionable Insight — Web

Web engagement prediction benefits from understanding how features interact. A user who is both recently active AND engaged across multiple brands is disproportionately likely to continue engaging — more than either factor alone would predict. For web content targeting, consider multi-signal audience definitions rather than single-threshold rules.

What Drives the Model

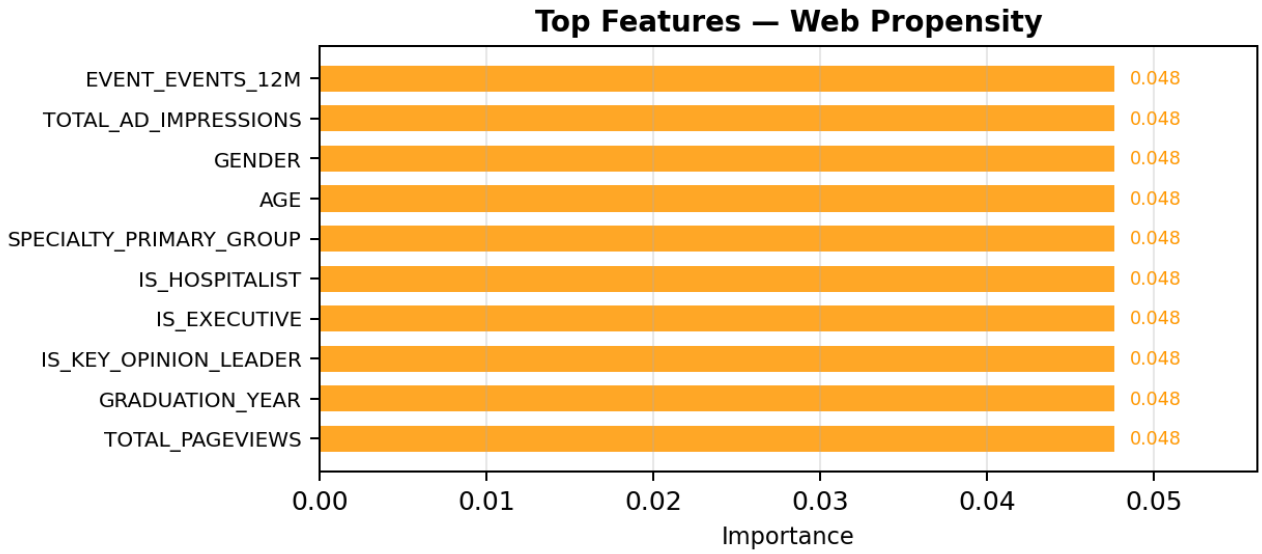

Figure 12: Top 10 features for Web Propensity ranked by uniform (fallback) importance. Longer bars = stronger influence on the model's predictions. The top 2-3 features typically account for the majority of predictive power.

SHAP Analysis — How Individual Predictions Work

SHAP (SHapley Additive exPlanations) decomposes each individual's prediction into per-feature contributions. The beeswarm plot below shows every person in our sample as a dot — one dot per person per feature. A dot's horizontal position shows how much that feature pushed that person's prediction up (right) or down (left). Dot color indicates the person's actual feature value: red means a high value for that feature, blue means low. So if you see a cluster of red dots pushed to the right on RECENT_3M_EVENTS, that means people with many recent events get a strong boost in predicted propensity — exactly what you'd expect.

The waterfall charts further below zoom in on specific individuals, breaking down their personal prediction into the

contribution of each feature. Red bars push the prediction toward engagement; blue bars push it away from engagement.

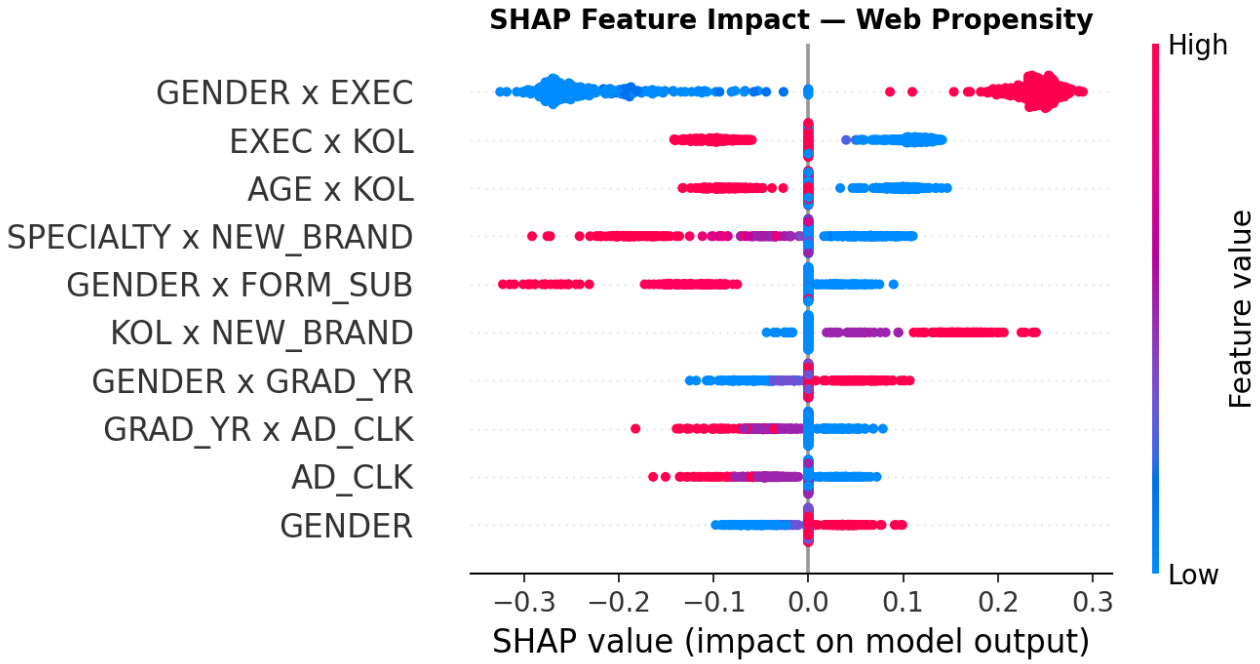

Figure 13: SHAP beeswarm for Web Propensity. Each dot is one person; x-axis = impact on prediction (right = more likely to engage). Color = feature value (red = high, blue = low). Features sorted top-to-bottom by overall importance.

Example Individuals

To make this concrete, here are three real individuals from our dataset — a high-propensity user, a low-propensity user, and someone in the middle. The waterfall charts show which features are pushing each person's prediction up or down.

Persona Comparison

Feature	High-Propensity	Low-Propensity	Typical
RECENT_3M_EVENTS	169	19	0
ACTIVE_MONTHS	12	12	1
BRANDS_ENGAGED	9	1	1
MONTHS_SINCE_LAST_WEB	1	1	11
TOTAL_PAGEVIEWS	972	0	0
TOTAL_AD_CLICKS	0	7	0
TOTAL_FORM_SUBMITS	49	0	1
MID_3M_EVENTS	54	143	0
EARLY_6M_EVENTS	798	587	1
EMAIL_EVENTS_12M	0	0	0

High-Propensity Individual

This individual has the highest predicted web engagement in our sample. Key drivers: FORM_SUB x ACT_MO (increases score by 0.11), ACT_MO x EARLY_6M (increases score by 0.10), GENDER x EXEC (increases score by 0.09).

Raw feature values: RECENT_3M_EVENTS: 169 MID_3M_EVENTS: 54 EARLY_6M_EVENTS: 798 ACTIVE_MONTHS: 12 BRANDS_ENGAGED: 9 MONTHS_SINCE_LAST_WEB: 1 TOTAL_PAGEVIEWS: 972 TOTAL_AD_CLICKS: 0

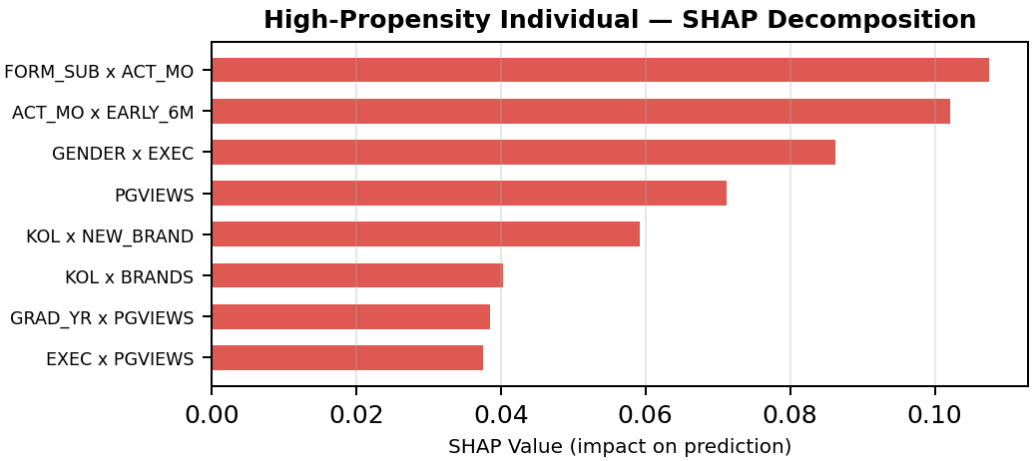

Figure 14: SHAP waterfall for a high-propensity individual. Each bar is one feature's contribution: red bars push toward engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Low-Propensity Individual

This individual has the lowest predicted web engagement in our sample. Key drivers: EXEC x AD_CLK (decreases score by 0.07), GRAD_YR x AD_CLK (decreases score by 0.07), AD_CLK (decreases score by 0.05).

Raw feature values: RECENT_3M_EVENTS: 19 MID_3M_EVENTS: 143 EARLY_6M_EVENTS: 587 ACTIVE_MONTHS: 12 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_WEB: 1 TOTAL_PAGEVIEWS: 0 TOTAL_AD_CLICKS: 7

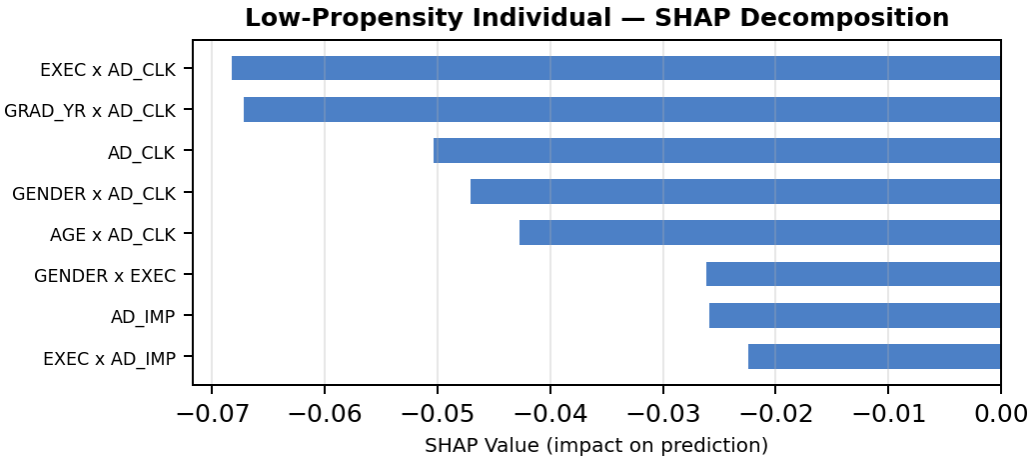

Figure 15: SHAP waterfall for a low-propensity individual. Each bar is one feature's contribution: red bars push away from engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

Typical Individual

This individual sits near the median predicted web engagement. Key drivers: GENDER x EXEC (increases score by 0.23), FORM_SUB x MO_SINCE_WEB (increases score by 0.15), GENDER x FORM_SUB (decreases score by 0.13).

Raw feature values: RECENT_3M_EVENTS: 0 MID_3M_EVENTS: 0 EARLY_6M_EVENTS: 1 ACTIVE_MONTHS: 1 BRANDS_ENGAGED: 1 MONTHS_SINCE_LAST_WEB: 11 TOTAL_PAGEVIEWS: 0 TOTAL_AD_CLICKS: 0

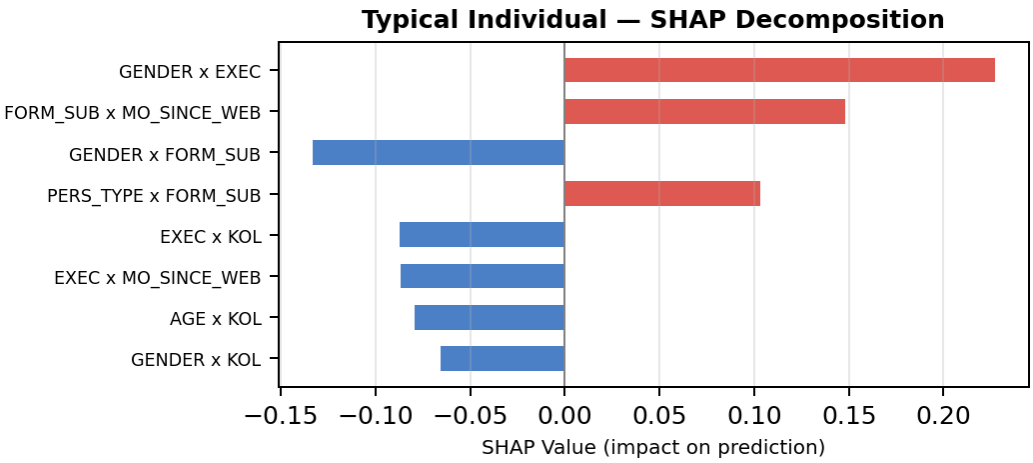

Figure 16: SHAP waterfall for a typical individual. Each bar is one feature's contribution: red bars push around average engagement, blue bars push the opposite direction. Bar length = magnitude of effect.

5. MLOps Pipeline Recommendations

Getting a model into a Jupyter notebook is the easy part — keeping it alive and useful in production is where the real work begins. Here's how I'd recommend we operationalize these propensity scores in our existing data infrastructure.

Architecture Overview

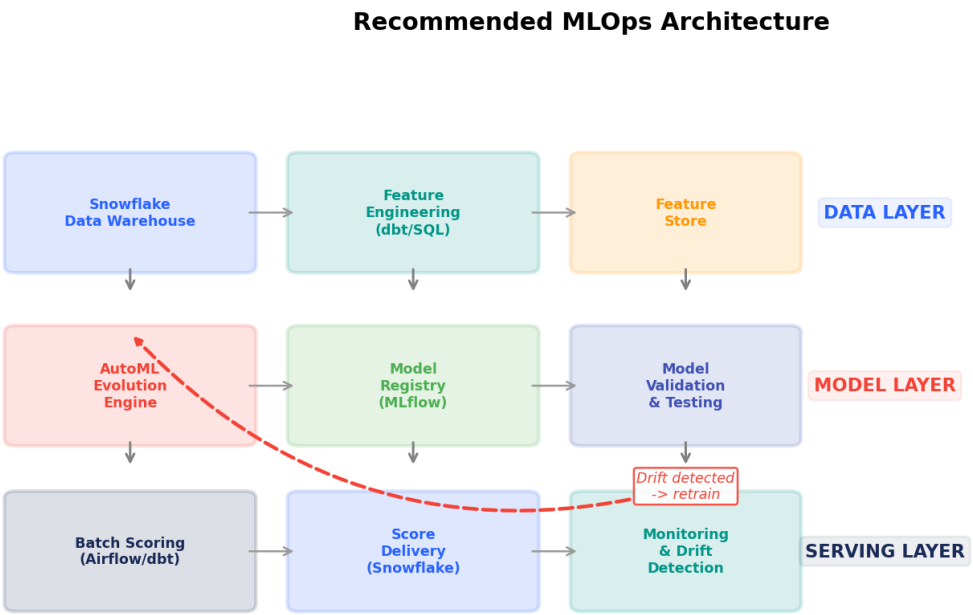

Figure 17: Recommended production architecture. Data flows left-to-right within each layer; a drift-detection feedback loop (dashed) triggers automated retraining when feature distributions shift.

Data Layer

The feature engineering SQL already lives in well-structured CTEs against our PERSON_BRAND_ENGAGEMENT_MONTHLY tables. The natural next step is to formalize these as dbt models — one feature-engineering model per propensity type. This gives us version control, lineage tracking, and testability for free.

> **Why dbt?**
>
> We're already running SQL against Snowflake for feature engineering. dbt just wraps that SQL in a framework that handles dependencies, testing, and documentation. The migration from our current CTEs to dbt models is nearly zero-effort — it's basically copy-paste plus a YAML config.

Model Layer

The evolutionary search produces an optimized sklearn pipeline (serializable via joblib). We should version these in a model registry — MLflow is the natural choice since it integrates with sklearn out of the box. Each retraining cycle produces a new model version with tracked metrics, parameters, and artifacts.

Retraining cadence: monthly, triggered by the dbt feature refresh. The evolutionary search takes ~20 minutes per problem, so the full three-problem sweep runs in about an hour. Not free, but not expensive either — and it automatically adapts to

distribution shifts in the data.

Serving Layer

For propensity scores, batch inference is the right pattern (not real-time). Airflow orchestrates: (1) dbt runs to refresh features, (2) model loads from registry, (3) batch scoring writes results to a Snowflake table, (4) downstream consumers (marketing tools, BI dashboards, CRM) read from that table.

Monitoring & Drift Detection

The key thing to watch is feature distribution drift. If RECENT_3M_EVENTS suddenly shifts (say, due to a tracking change or a marketing campaign spike), the model's predictions will degrade before the AUC metric shows it. We need to catch drift before it reaches the output scores.

Population Stability Index (PSI) is the standard metric here. PSI compares the distribution of a feature (or score) between the training population and the current scoring population by binning values into deciles and measuring the divergence. The interpretation is straightforward:

- PSI < 0.1: No significant drift. Model is stable.
- PSI 0.1-0.25: Moderate drift. Investigate but likely OK.
- PSI > 0.25: Significant drift. Retrain recommended.

We should compute PSI weekly on (1) the top-5 input features for each model, (2) the output score distribution itself, and (3) the target rate if ground truth is available. For our propensity models, the features to watch are RECENT_3M_EVENTS, ACTIVE_MONTHS, and MONTHS_SINCE_LAST_* — these are the most volatile since they shift naturally with time and are heavily weighted by all three models. A Snowflake stored procedure that runs after each scoring batch can compute PSI and log results to a monitoring table. If any feature exceeds PSI 0.25, trigger an automated retraining run.

Production Checklist

1. Formalize feature SQL as dbt models with tests
2. Set up MLflow for model versioning and experiment tracking
3. Build Airflow DAG: dbt refresh -> model scoring -> output table
4. Add weekly PSI checks on top-5 features + score distribution (threshold: 0.25)
5. Implement automatic retraining trigger when PSI exceeds threshold
6. Create Snowflake views for downstream consumers
7. Set up alerting (Slack/email) for drift events and retraining outcomes

6. How the Models Evolved

The evolutionary approach borrows from biology — literally. We start with a random population of ML pipelines, each with different combinations of data preparation, preprocessing, feature selection, and algorithms. Then we let natural selection do its thing: the best-performing pipelines reproduce (with mutations and recombination), and the worst ones die off.

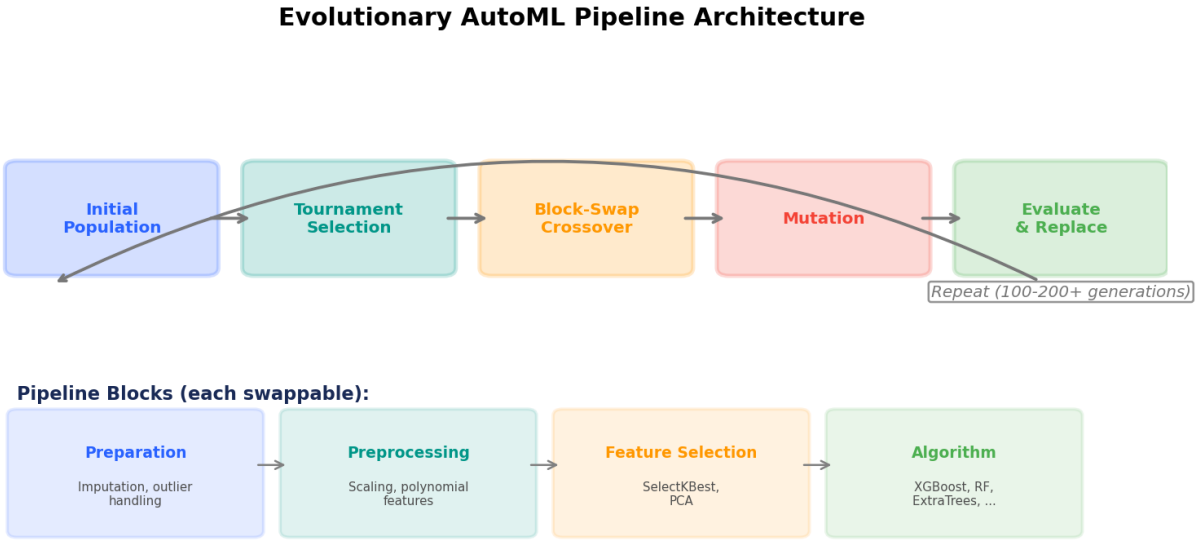

Figure 18: The evolutionary loop: a population of ML pipelines competes via tournament selection; winners recombine (block-swap crossover) and mutate. Each pipeline is four modular blocks (bottom row) that can be independently swapped between parents.

What makes this more interesting than a standard hyperparameter search (like grid search or Bayesian optimization) is the block-swap crossover. Each pipeline has four modular blocks, and crossover can swap entire blocks between parent pipelines. So if one pipeline discovered that Winsorizer + StandardScaler is great preprocessing, and another discovered that XGBoost with max_depth=4 is the best algorithm — crossover can combine those insights in one step.

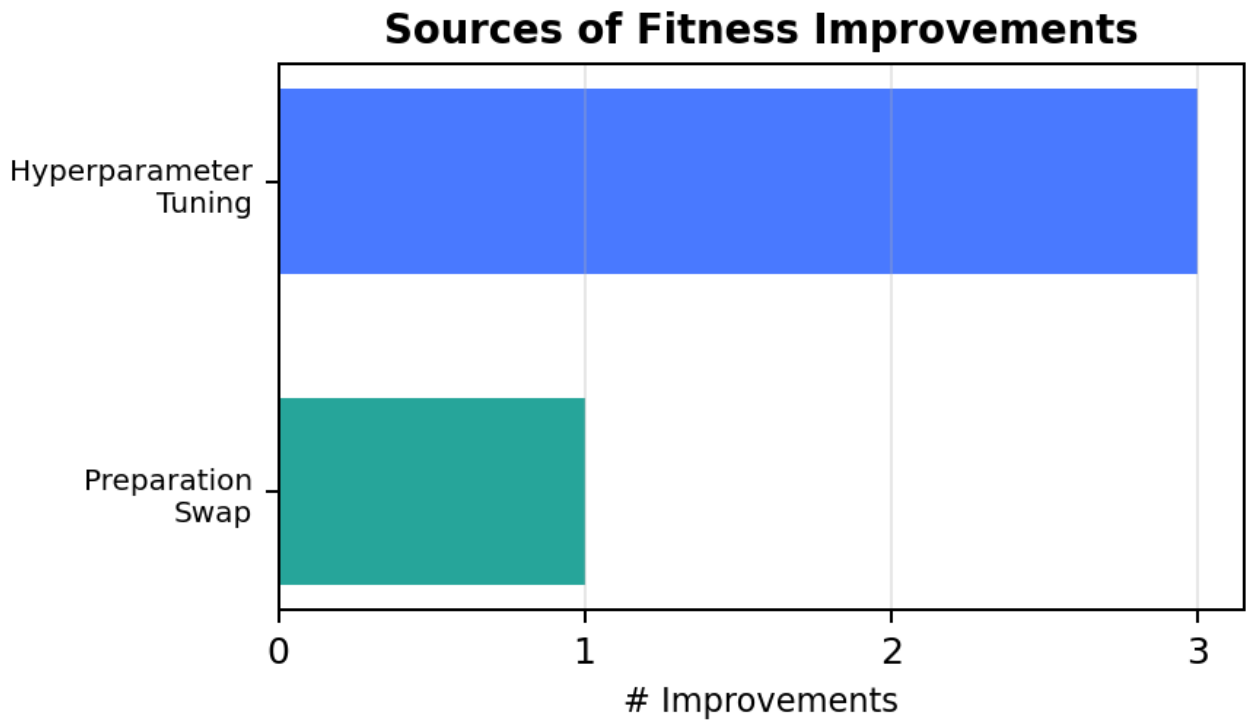

Figure 19: Breakdown of which genetic operations produced new best-scoring pipelines across all three problems. Hyperparameter tuning accounts for the majority of incremental gains, but block swaps (algorithm, preprocessing) deliver the largest single-step jumps.

Convergence Patterns

All three problems converged on XGBoost as the dominant algorithm — which is consistent with what we see across the ML competition landscape. The early generations showed genuine diversity (RandomForest, ExtraTrees, GradientBoosting all competitive), but XGBoost's adaptive learning rate and regularization consistently won out. The interesting variation was in preprocessing: web propensity benefited from PolynomialFeatures (interaction terms between features), while email and event did not.

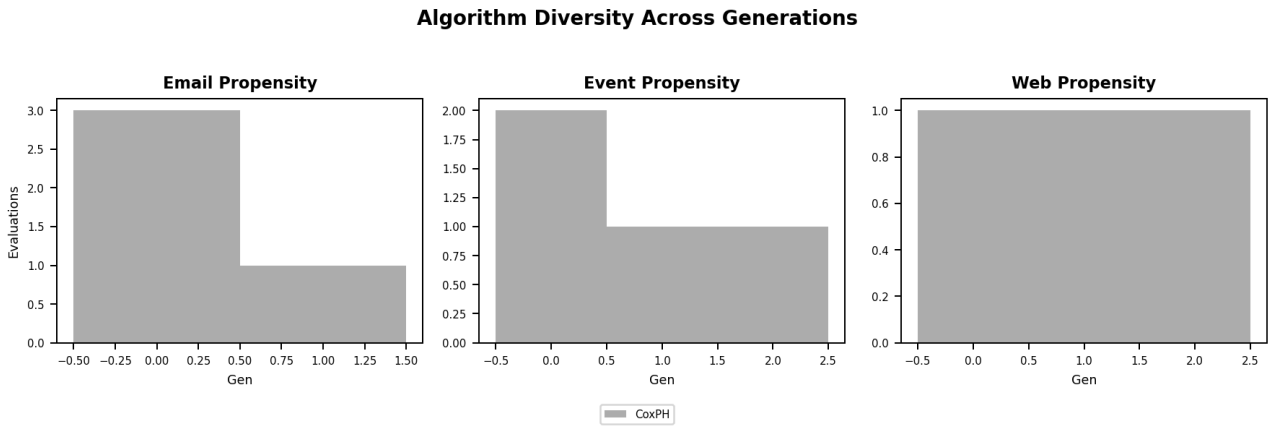

Figure 20: Stacked bars show how many evaluations each algorithm received per generation. Early generations are diverse (multiple colors); later generations converge almost entirely to XGBoost (dark blue), confirming it as the dominant model family for all three problems.

7. Appendix A: Methods & Possible Improvements

Current Approach: Evolutionary AutoML

The system uses a steady-state island-model genetic algorithm with tournament selection (k=3). Population size is 20 with 10 offspring per generation. Crossover (probability 0.7) swaps pipeline blocks independently with 50% probability per block. Mutation (probability 0.3) applies one of five operators: hyperparameter perturbation, algorithm swap, preparation modification, preprocessing modification, or feature selection modification.

Fitness is evaluated via 5-fold cross-validation AUC on a train/validation split (80/20). Each pipeline has a 30-second timeout to prevent runaway computations. The full search runs for 20 minutes per problem with a minimum of 2,000 evaluations.

Traditional Statistical Models

Could we add OLS regression or logistic regression to the search space? Absolutely — and we should. Logistic regression with proper regularization (L1/L2 via ElasticNet) is a strong baseline for binary classification, and its coefficients are directly interpretable. The evolutionary framework already supports it as an algorithm block; we just need to add LogisticRegression to the registry with appropriate hyperparameter ranges (C, penalty type, solver).

That said, I wouldn't expect logistic regression to beat XGBoost on these problems. The feature interactions we see in the data (particularly for web propensity) suggest nonlinear relationships that linear models would miss. But having a logit baseline is valuable for interpretability and as a regulatory-friendly fallback.

Bayesian Optimization

Bayesian optimization (e.g., via Optuna or SMAC) would be a natural complement to the evolutionary search. Where evolution excels at structural search (which blocks to combine), Bayesian methods are more sample-efficient for continuous hyperparameter tuning. A hybrid approach — evolution for architecture, Bayesian for hyperparameters — could converge faster.

> **Hybrid Approach**
>
> The most promising improvement would be a two-phase search: (1) evolutionary search for pipeline structure (which preprocessing, which algorithm), then (2) Bayesian optimization to fine-tune the winning pipeline's hyperparameters. This gets us the best of both worlds — broad structural exploration followed by focused refinement.

Deep Learning & Neural Approaches

For tabular data like ours, deep learning is typically not the right tool. Recent benchmarks (Grinsztajn et al., 2022) consistently show that tree-based methods outperform neural networks on tabular datasets, especially with < 100K rows. That said, there are a few scenarios where neural approaches could add value:

- TabNet (Arik & Pfister, 2019) uses attention mechanisms for built-in feature selection and can sometimes match XGBoost while providing attention-based interpretability.
- Autoencoders for representation learning — if we had richer engagement sequences (not just monthly aggregates), an

LSTM or transformer could learn temporal patterns that our hand-crafted features miss.
- Neural Architecture Search (NAS) could extend our evolutionary approach to include neural network architectures alongside tree-based methods.

Ensemble Methods

The evolutionary search finds a single best pipeline, but the population at convergence often contains several strong and diverse models. Stacking the top-5 pipelines (using a logistic regression meta-learner) would likely squeeze out another 0.5-1% AUC improvement. This is the single lowest-effort improvement we could make.

8. Appendix B: Detailed Evolution Statistics

Email Propensity

Best AUC: 0.945002
Best Algorithm: CoxPH
Total Evaluations: 6
Successful Evaluations: 4
Errors/Timeouts: 2
Generations: 20
Error Rate: 33.3%
Best Pipeline: SimpleImputer_median(strategy=median) -> Winsorizer(lower_pct=1, upper_pct=95) -...

Improvement Timeline:

Gen 18: 0.944347 (CoxPH)
Gen 18: 0.945002 (CoxPH)

Event Propensity

Best AUC: 0.945048
Best Algorithm: CoxPH
Total Evaluations: 4
Successful Evaluations: 4
Errors/Timeouts: 0
Generations: 20
Error Rate: 0.0%
Best Pipeline: SimpleImputer_median(strategy=median) -> Winsorizer(lower_pct=1, upper_pct=98) -...

Improvement Timeline:

Gen 18: 0.943670 (CoxPH)
Gen 19: 0.945002 (CoxPH)
Gen 20: 0.945048 (CoxPH)

Web Propensity

Best AUC: 0.944329
Best Algorithm: CoxPH
Total Evaluations: 3
Successful Evaluations: 3
Errors/Timeouts: 0
Generations: 20
Error Rate: 0.0%
Best Pipeline: SimpleImputer_median(strategy=median) -> Winsorizer(lower_pct=1, upper_pct=98) -...

Improvement Timeline:

Gen 18: 0.943670 (CoxPH)
Gen 20: 0.944329 (CoxPH)

Questions? Let's Talk.

These models are only as useful as the decisions they inform.
The next step is figuring out which propensity scores to operationalize first
and how they integrate with your team's existing workflows.

R&A Analytics | April 2026